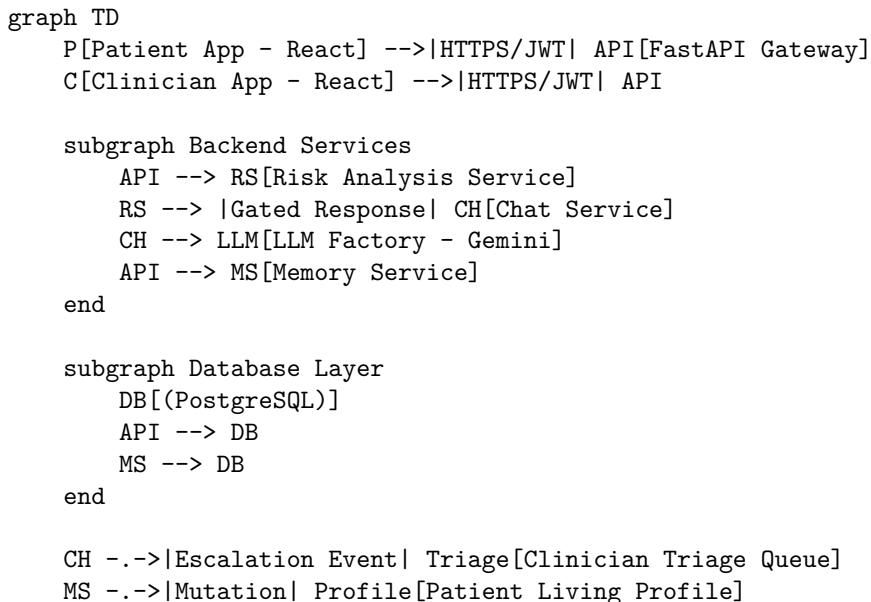


Technical Brief: Nightingale Medical Messaging System

1. Architecture Overview

The Nightingale system is built as a **Secure, AI-Gated Healthcare Messaging Platform**. It employs a “Safety-First” architecture where every interaction is validated for medical risk before reaching the patient.



Key Architectural Pillars:

- **Risk-Gated Pipeline:** Every incoming message passes through `RiskAnalysisService` before any AI advice is generated. If “Med” or “High” risk is detected, the AI is short-circuited.
- **Background Extraction:** Memory mutation (extracting facts like medications) happens asynchronously via background tasks to ensure sub-second response latency for the patient.
- **Clinician-in-the-Loop:** The “Escalation Loop” ensures that AI never operates in a vacuum for critical cases.

2. Data Schema

The schema is designed for **Medical Provenance**. Every extracted fact must point back to the specific message that generated it.

```

erDiagram
    USER ||--o{ CONVERSATION : owns
    USER {
        int id
        string email
        string role "patient | clinician"
    }

    CONVERSATION ||--o{ MESSAGE : contains
    MESSAGE ||--o{ CITATION : supports
    MESSAGE {
        int id
        string content_redacted
        string role "user | assistant | clinician"
        string risk_level
    }

    MESSAGE ||--o{ ESCALATION : triggers
    ESCALATION {
        int id
        string triage_summary
        string status "pending | resolved"
    }

    USER ||--o{ MEDICAL_MEMORY : has
    MEDICAL_MEMORY {
        int id
        string value "Advil"
        string status "active | stopped"
        int message_id "Provenance Link"
    }

    MESSAGE ||--o{ AUDIO_TRANSCRIPT : "Voice Readiness"
    AUDIO_TRANSCRIPT {
        int id
        string storage_path
        string status "processed | failed"
    }

```

3. First-Principles Thinking & Assumptions

1. **Safety Over Fluency:** It is better for the AI to refuse to answer a high-risk question than to give a hallucinated diagnosis.
2. **Provenance is Trust:** In a clinical setting, a fact (e.g., “Allergic to

- Penicillin") is useless without knowing *who* said it and *when*.
3. **Privacy by Design:** PHI (Patient Health Information) like NRIC or Phone numbers should never enter the LLM's context window. Redaction happens at the edge (Backend API).
 4. **Empower, Don't Replace:** The AI is a triage assistant and educator, not a doctor. Every interaction ends with a nudge toward professional consultation.
-

4. Trade-offs & Scope

What was prioritized:

- **The “Safety Stop”:** Ensuring high-risk symptoms (like chest pain) immediately disable AI and alert the clinic.
- **Living Memory:** Implementing the logic that handles “I take X” vs “Actually I stopped X” correctly within a single session.
- **Clinician Ground Truth:** Ensuring that when a nurse replies, that message is marked as “Verified” and used to ground future AI responses.

What was cut (Out of Scope):

- **Complex Clinician Roster:** In this MVP, all clinicians see a global triage queue. In production, this would be scoped to specific clinics/specialties.
 - **Full E2E Encryption:** While TLS is used, the backend still processes text to perform redaction and extraction. A future iteration might use Trusted Execution Environments (TEEs).
-

5. VoiceAI Strategy

Nightingale is “Voice-Ready” by design. Here is how audio slots into the current schema:

1. **Schema Support:** The `Message` table includes an optional `audio_id` and `transcript_id`.
2. **Ingestion Flow:**
 - Patient records audio in the Frontend.
 - The `.wav` file is uploaded to secure object storage (e.g., S3/GCS).
 - An asynchronous STT (Speech-to-Text) service processes the audio.
 - The resulting text is fed into the same `RiskAnalysisService` and `MemoryService` pipeline.
3. **Grounding:** The `MEDICAL_MEMORY` provenance link would point to the transcript, and the UI would allow the clinician to “Play Original Audio” to verify intent/tone (e.g., slurred speech as a symptom).